

3. Theory solvers, combinations and extensions

Prof. Roberto Amadini

Department of Computer Science and Engineering, University of Bologna, Italy

Combinatorial Decision Making and Optimization

2nd cycle degree programme in Artificial Intelligence

University of Bologna, Academic Year 2024/25



Theory solvers

- So far, we introduced the basics of SMT solving without focusing much on background **theories** and their **solvers**
 - Eager vs lazy approaches
- In its simplest form, a **\mathcal{T} -solver** takes as input a **conjunction** of \mathcal{T} -literals μ and **decides** whether μ is **\mathcal{T} -satisfiable**
- We can see a SMT solver as a “collection” of theory solvers
- What are the crucial **features** for a \mathcal{T} solver?

Theory solvers

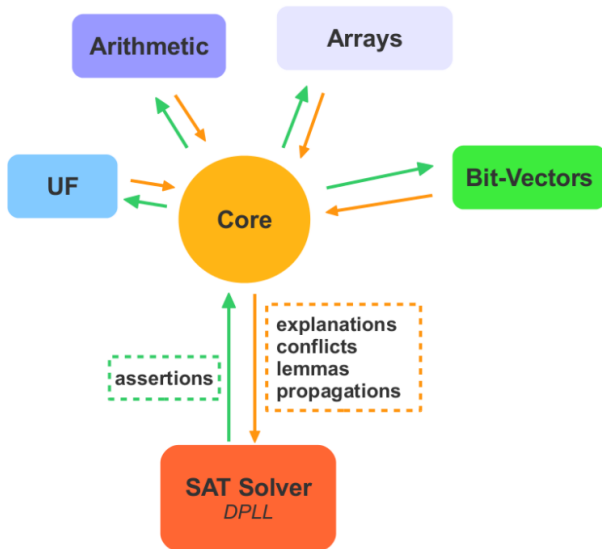
- **Early pruning**: invoke \mathcal{T} -solver on **partial** Boolean assignments, especially during the early stages of the search
- **Incrementality**: when a new constraint is **added**, possibly avoid recomputing everything from scratch
- **Backtrackability**: support cheap (stack-based) **removal** of constraints when exploring the search tree without “resetting” the internal state

- **Literal deduction**: \mathcal{T} -solver can perform **deductions** of literals not yet assigned in the input formula (**\mathcal{T} -propagation**)
- **Explanation generation**: when a conflict involving a literal ℓ is found, is necessary to get a (possibly short) **explanation** $\ell_1 \wedge \dots \wedge \ell_n \rightarrow \ell$ to perform **conflict analysis** and **backjumping**

What theories?

- Uninterpreted functions (EUF)
- Arithmetic
 - LIA, LRA, LIRA, ...
- Arrays
- Bit-vectors
- Strings
- ...


What theories?



- A \mathcal{T}_{EUf} formula can be decided in polynomial time with congruence closure procedure:
 - Add fresh c and replace each $p(t_1, \dots, t_k)$ with $f_p(t_1, \dots, t_k) = c$
 - Partition formula into equalities (E) and disequalities (D)
 - Compute the congruence closure \equiv_E of E , i.e., the smallest equivalence relation over the terms of E such that:
 - $t_1 = t_2 \in E \implies t_1 \equiv_E t_2$
 - For each $f(s_1, \dots, s_k)$, $f(t_1, \dots, t_k)$ occurring in E , if $s_i \equiv_E t_i$ for each $i \in \{1, \dots, k\}$ then $f(s_1, \dots, s_k) \equiv_E f(t_1, \dots, t_k)$ (congruence property)
 - The formula is satisfiable iff for each $t_1 \neq t_2 \in D$ we have $t_1 \not\equiv_E t_2$
- Standard algorithms use a DAG to represent functions applications, and union-find (a.k.a. merge-find or disjoint-set) for the classes of \equiv_E

EUF theory

Example*: $\phi \equiv f(a, b) = a \wedge f(f(a, b), b) = c \wedge g(a) \neq g(c)$



(a) DAG for ϕ -terms. (b) E-graph: equivalences are the equalities in ϕ .
(c) $f(f(a, b), b) \equiv_E f(a, b)$ because $f(a, b) \equiv_E a$. (d) $g(a) \equiv_E g(c)$
because $a \equiv_E c$. Since $g(a) \neq g(c)$ and $g(a) \equiv_E g(c)$, ϕ is **unsatisfiable**

- Consider **LRA** = **L**inear **R**eal **A**rithmetic theory, having signature $\Sigma_{LRA} = (\mathbb{Q}, +, -, *, \leq)$ and **linear multiplications** only
- We could decide LRA-literals with **Fourier-Motzkin elimination**
 - Replace $t_1 \neq t_2$ with $t_1 < t_2 \vee t_2 < t_1$, and $t_1 \leq t_2$ with $t_1 < t_2 \vee t_1 = t_2$ (case splitting)
 - Eliminate equalities and apply Fourier-Motzkin elimination to all variables to determine its satisfiability
 - https://en.wikipedia.org/wiki/Fourier-Motzkin_elimination
- Not practical for large set of constraints, **simplex method** preferable

LIA theory

- Consider **LIA** = **L**inear **I**nteger **A**rithmetic theory, having signature $\Sigma_{LIA} = (\mathbb{Z}, +, -, *, \leq)$ and **linear multiplications** only
 - if not linear, undecidable (*Peano arithmetic*)
 - if fully quantified, *Presburger arithmetic*
 - if quantifier-free, different decision procedures exist
- As for LRA, we can apply methods like Fourier-Motzkin, but Simplex + **branch & bound/cut** generally better
- Methods exist also for **LIRA** = integer + real arithmetic and **NLA** = nonlinear arithmetic
 - E.g., <https://microsoft.github.io/z3guide/docs/theories/Arithmetic/>

Difference logic

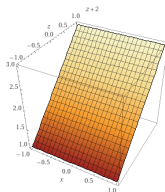
- Consider now **DL** = **D**ifference **L**ogic theory, having atomic formulas of the form $x - y \leq k$ with x, y variables and k constant
 - Constraints $x - y \bowtie k$ with $\bowtie \in \{=, \neq, <, \geq, >\}$ can be **rewritten**

- E.g. if $x, y \in \mathbb{Z}$:

$$x - y > 5 \wedge x = z + 2 \implies$$

$$x - y \geq 6 \wedge x - z \leq 2 \wedge x - z \geq 2 \implies$$

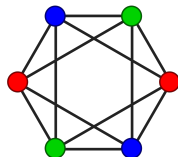
$$y - x \leq -6 \wedge x - z \leq 2 \wedge z - x \leq -2$$



- Unary constraints $x \leq k$ can be rewritten into $x + z_0 \leq k$ by enforcing $z_0 = 0$ in any satisfying assignment


Difference logic and k -coloring

- If we allow \neq and the domain is \mathbb{Z} , deciding satisfiability of DL formulas is **NP-hard**, e.g., it gets “as hard as” **k -coloring problem**
 - If we have k colors available, can we color a graph s.t. **adjacent** nodes have **different** colors? If $k \geq 3$ the problem is NP-hard
- Formally, given graph (V, E) and $k \in \mathbb{N}$, does it exist a function $c : V \rightarrow \{1, \dots, k\}$ s.t. for each $(i, j) \in E$ we have $c(i) \neq c(j)$?
- **Any** k -coloring instance can be mapped to a DL formula with $|V|$ variables, $|E|$ disequalities $x_i \neq x_j$ for each $(i, j) \in E$ and $2|V|$ disequalities $1 \leq x_i \leq k$
 - If we can decide the DL formula in polynomial time, we can solve **any** problem of NP in polynomial time



Difference logic as graph problem

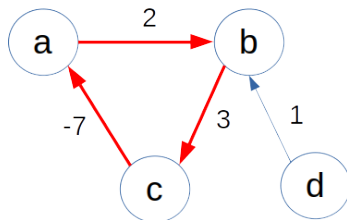
- From DL literals in Φ we can get a **directed weighted graph** \mathcal{G}_Φ with:
 - a node for each **variable** occurring in Φ
 - a weighted edge $x \xrightarrow{k} y$ for each $x - y \leq k \in \Phi$
- E.g., if $\Phi = \{y - x \leq -6, x - z \leq 2, z - x \leq -2\}$ then \mathcal{G}_Φ is:



- Theorem:** Φ is inconsistent $\iff \mathcal{G}_\Phi$ has a **negative cycle**

Difference logic as graph problem

- Example: let $\Phi = \{a - b \leq 2, b - c \leq 3, c - a \leq -7, d - b \leq 1\}$.
The \mathcal{G}_Φ graph is:



- Negative loop $a \xrightarrow{2} b \xrightarrow{3} c \xrightarrow{-7} a$ (total weight -2): Φ **inconsistent**
 - $a \geq c + 7$ conflicts with $a \leq b + 2 \leq (c + 3) + 2 = c + 5$

Difference logic as graph problem

- Negative loops can be detected with **Bellman-Ford** in $O(|V||E|)$ by adding to V a **source** vertex x_0 and an edge $x_0 \xrightarrow{0} x$ for each $x \in E$
 - https://en.wikipedia.org/wiki/Bellman%E2%80%93Ford_algorithm
 - Other more efficient variant exists
- Negative loops denotes **inconsistency explanations**
 - Not minimal in general
- Theory **propagations** computed from consistent graphs: if there is a **path** between x and y with total weight k , we can **deduce** $x - y \leq k$
 - If $x \xrightarrow{k_1} x_1 \xrightarrow{k_2} x_2 \xrightarrow{k_3} \dots \xrightarrow{k_n} y$ the total weight is $k = \sum_{i=1}^n k_i$ and $x - x_1 \leq k_1, x_1 - x_2 \leq k_2, \dots, x_n - y \leq k_n$ hence $(x - x_1) + (x_1 - x_2) + \dots + (x_n - y) \leq \sum_{i=1}^n k_i = k$ thus $x + (-x_1 + x_1) + \dots + (-x_n + x_n) + y \leq k$, i.e., $x - y \leq k$

Other theories

- **Bit-vectors**: typically BV formulas are first **simplified**, and then encoded into SAT formulas (**bit-blasting**)
- Arrays: typically flattening of terms + congruence closure + lazy axioms instantiation + optimizations
- (Multi)-sets
- Strings
- Floating points
- ...

Combining Theories

Need for combination

- So far we considered theories individually. But often SMT formulas contain atoms from (very) **different theories**
- In particular **software verification** applications can generate constraints over several data types
 - integers, floating points, bit-vectors, arrays, strings, ...
- E.g., formula $a = b + 2 \wedge A = \text{write}(B, a + 1, 4) \wedge (f(a) \vee g(b + 1))$ involves theory of linear arithmetic, arrays, and EUF
- Given **\mathcal{T}_i -solvers** for theories $\mathcal{T}_1, \dots, \mathcal{T}_n$, can we **combine** them to get a solver for $\bigcup_i \mathcal{T}_i$?

Example

- Consider formula $f(f(x) - f(y)) = a \wedge f(0) = a + 2 \wedge x = y$
 - Two theories involved: **EUF** and linear arithmetic (**LA**)
- 1st step: purification**. Each literal must belong to **only one** theory
 - Fresh** constants needed: e_1, e_2, \dots, e_5
- Purified formula:
$$\underbrace{e_1 = e_2 - e_3}_{\text{LA}}, \underbrace{f(e_1) = a, e_2 = f(x), e_3 = f(y)}_{\text{EU F}},$$
$$\underbrace{e_4 = 0, e_5 = a + 2}_{\text{LA}}, \underbrace{f(e_4) = e_5, x = y}_{\text{EU F}}$$
- In this way EUF and LA solvers only **share** a, e_1, \dots, e_5
- To **merge** the corresponding models, solvers must **agree** on equalities between shared constants, a.k.a. **interface equalities**

Example

- 2nd step: satisfiability **check** and equalities **exchange**
- Purified formula: $\underbrace{e_1 = e_2 - e_3}_{LA}, \underbrace{f(e_1) = a, e_2 = f(x), e_3 = f(y)}_{EUF},$
 $\underbrace{e_4 = 0, e_5 = a + 2}_{LA}, \underbrace{f(e_4) = e_5, x = y}_{EUF}$
- Both EUF-solver and LA-solver say **SAT**
- **EUF solver** deduces that $\{x = y, f(x) = e_2, f(y) = e_3\} \models e_2 = e_3$ and **sends** the literal to the LA solver

Example

- 2nd step: satisfiability **check** and equalities **exchange**
- Purified formula: $\underbrace{e_1 = e_2 - e_3}_{LA}, \underbrace{f(e_1) = a, e_2 = f(x), e_3 = f(y)}_{EUF},$
 $\underbrace{e_4 = 0, e_5 = a + 2, e_2 = e_3}_{LA}, \underbrace{f(e_4) = e_5, x = y}_{EUF}$
- Both EUF-solver and LA-solver say **SAT**
- **EUF solver** deduces that $\{x = y, f(x) = e_2, f(y) = e_3\} \models e_2 = e_3$ and **sends** the literal to the LA solver

Example

- 2nd step: satisfiability check and equalities exchange
- Purified formula: $\underbrace{e_1 = e_2 - e_3}_{LA}, \underbrace{f(e_1) = a, e_2 = f(x), e_3 = f(y)}_{EUF},$
 $\underbrace{e_4 = 0, e_5 = a + 2, e_2 = e_3}_{LA}, \underbrace{f(e_4) = e_5, x = y}_{EUF}$
- Both EUF-solver and LA-solver say SAT
- LA solver deduces that $\{e_2 - e_3 = e_1, e_4 = 0, e_2 = e_3\} \models e_1 = e_4$ and sends the literal to the EUF solver

Example

- 2nd step: satisfiability check and equalities exchange
- Purified formula: $\underbrace{e_1 = e_2 - e_3}_{LA}, \underbrace{f(e_1) = a, e_2 = f(x), e_3 = f(y)}_{EUF},$
 $\underbrace{e_4 = 0, e_5 = a + 2, e_2 = e_3}_{LA}, \underbrace{f(e_4) = e_5, x = y, e_1 = e_4}_{EUF}$
- Both EUF-solver and LA-solver say SAT
- LA solver deduces that $\{e_2 - e_3 = e_1, e_4 = 0, e_2 = e_3\} \models e_1 = e_4$ and sends the literal to the EUF solver

Example

- 2nd step: satisfiability **check** and equalities **exchange**
- Purified formula: $\underbrace{e_1 = e_2 - e_3}_{LA}, \underbrace{f(e_1) = a, e_2 = f(x), e_3 = f(y)}_{EUF},$
 $\underbrace{e_4 = 0, e_5 = a + 2, e_2 = e_3}_{LA}, \underbrace{f(e_4) = e_5, x = y, e_1 = e_4}_{EUF}$
- Both EUF-solver and LA-solver say **SAT**
- **EUF solver** deduces that $\{f(e_1) = a, f(e_4) = e_5, e_1 = e_4\} \models a = e_5$
and sends the literal to the LA solver

Example

- 2nd step: satisfiability **check** and equalities **exchange**
- Purified formula: $\underbrace{e_1 = e_2 - e_3}_{LA}, \underbrace{f(e_1) = a, e_2 = f(x), e_3 = f(y)}_{EUF},$
 $\underbrace{e_4 = 0, e_5 = a + 2, e_2 = e_3, a = e_5}_{LA}, \underbrace{f(e_4) = e_5, x = y, e_1 = e_4}_{EUF}$
- Both EUF-solver and LA-solver say **SAT**
- **EUF solver** deduces that $\{f(e_1) = a, f(e_4) = e_5, e_1 = e_4\} \models a = e_5$ and sends the literal to the LA solver

Example

- 2nd step: satisfiability **check** and equalities **exchange**
- Purified formula: $\underbrace{e_1 = e_2 - e_3}_{LA}, \underbrace{f(e_1) = a, e_2 = f(x), e_3 = f(y)}_{EUF},$
 $\underbrace{e_4 = 0, e_5 = a + 2, e_2 = e_3, a = e_5}_{LA}, \underbrace{f(e_4) = e_5, x = y, e_1 = e_4}_{EUF}$
- EUF-solver say **SAT**...
- ...but LA-solver say **UNSAT**: $\{e_5 = a + 2, a = e_5\} \models \perp$
- Hence the original formula is **UNSAT**

Nelson-Oppen procedure (convex case)

- Let Σ_1, Σ_2 be signatures and $\mathcal{T}_1, \mathcal{T}_2$ their theories. If \mathcal{T}_1 and \mathcal{T}_2 are:
 - **signature-disjoint**
 - $\Sigma_1 \cap \Sigma_2 = \emptyset$
 - **stably-infinite**
 - Σ -theory \mathcal{T} of sort σ is stably infinite if every \mathcal{T} -satisfiable Σ -formula has a model interpreting σ as an **infinite set**
 - **convex**
 - For each set of \mathcal{T}_i -literals S , $S \models_{\mathcal{T}_i} (a_1 = b_1 \vee \dots \vee a_n = b_n)$ must imply that $S \models a_k = b_k$ for some $k \in \{1, \dots, n\}$
- then we can check the $(\mathcal{T}_1 \cup \mathcal{T}_2)$ -satisfiability with the **deterministic Nelson-Oppen** algorithm

Nelson-Oppen procedure (convex case)

Let S be a $(T_1 \cup T_2)$ -formula and E the set of **interface equalities** between S_1 and S_2 . **Deterministic Nelson-Oppen** steps:

1. **Purify** S and **split** it into S_1 and S_2
 - S_i contains \mathcal{T}_i -literals only
2. If $S_1 \models_{\mathcal{T}_1} \perp$, then return **UNSAT**
3. If $S_2 \models_{\mathcal{T}_2} \perp$, then return **UNSAT**
4. If $S_1 \models_{\mathcal{T}_1} (e = e')$ with $(e = e') \in E - S_2$, then $S_2 \leftarrow S_2 \cup \{(e = e')\}$ and go to 3
5. If $S_2 \models_{\mathcal{T}_2} (e = e')$ with $(e = e') \in E - S_1$, then $S_1 \leftarrow S_1 \cup \{(e = e')\}$ and go to 2
6. return **SAT**

Nelson-Oppen procedure (non-convex case)

- Why we needed convex theories? Consider the following formula:

$$1 \leq x \leq 2$$

$$f(1) = a$$

$$f(x) = b$$

$$a = b + 2$$

$$f(2) = f(1) + 3$$

involving linear integer arithmetic (**LIA**) and **EUF** theories

- Let's purify the formula by introducing the interface equalities:

$$e_1 = 1, \quad e_2 = 2, \quad e_3 = f(e_2), \quad e_4 = f(e_1), \quad e_3 = e_4 + 3$$

Nelson-Oppen procedure (non-convex case)

- Now let's **check** satisfiability and **exchange** entailed equalities:

<i>LIA</i>	<i>EUF</i>
$1 \leq x$	$f(e_1) = a$
$x \leq 2$	$f(x) = b$
$e_1 = 1$	$f(e_2) = e_3$
$a = b + 2$	$f(e_1) = e_4$
$e_2 = 2$	
$e_3 = e_4 + 3$	

- Both EUF-solver and LIA-solver say **SAT**
- EUF solver** deduces that $\{f(e_1) = a, f(e_1) = e_4\} \models a = e_4$ and sends the literal to the LA solver

Nelson-Oppen procedure (non-convex case)

- Now let's **check** satisfiability and **exchange** entailed equalities:

<i>LIA</i>	<i>EUF</i>
$1 \leq x$	$f(e_1) = a$
$x \leq 2$	$f(x) = b$
$e_1 = 1$	$f(e_2) = e_3$
$a = b + 2$	$f(e_1) = e_4$
$e_2 = 2$	
$e_3 = e_4 + 3$	
$a = e_4$	

- Both EUF-solver and LIA-solver say **SAT**
- EUF solver** deduces that $\{f(e_1) = a, f(e_1) = e_4\} \models a = e_4$ and sends the literal to the LA solver

Nelson-Oppen procedure (non-convex case)

- Now let's **check** satisfiability and **exchange** entailed equalities:

<i>LIA</i>	<i>EUF</i>
$1 \leq x$	$f(e_1) = a$
$x \leq 2$	$f(x) = b$
$e_1 = 1$	$f(e_2) = e_3$
$a = b + 2$	$f(e_1) = e_4$
$e_2 = 2$	
$e_3 = e_4 + 3$	
$a = e_4$	

- Both EUF-solver and LIA-solver say **SAT**
- EUF and LIA theories **cannot deduce** any other interface equality
 - ...but LIA solver could deduce $x = e_1 \vee x = e_2$

Nelson-Oppen procedure (non-convex case)

- Now let's **check** satisfiability and **exchange** entailed equalities:

<i>LIA</i>	<i>EUF</i>
$1 \leq x$	$f(e_1) = a$
$x \leq 2$	$f(x) = b$
$e_1 = 1$	$f(e_2) = e_3$
$a = b + 2$	$f(e_1) = e_4$
$e_2 = 2$	
$e_3 = e_4 + 3$	
$a = e_4$	

- If $x = e_1$, EUF would deduce $a = b$: **UNSAT**
- If $x = e_2$, EUF would deduce $b = e_3$: **UNSAT**

Nelson-Oppen procedure (non-convex case)

- Now let's **check** satisfiability and **exchange** entailed equalities:

<i>LIA</i>	<i>EUF</i>
$1 \leq x$	$f(e_1) = a$
$x \leq 2$	$f(x) = b$
$e_1 = 1$	$f(e_2) = e_3$
$a = b + 2$	$f(e_1) = e_4$
$e_2 = 2$	
$e_3 = e_4 + 3$	
$a = e_4$	

- Hence, $x = e_1 \vee x = e_2$ is false and the original formula **UNSAT**
- ...But we can't infer this with **deterministic** Nelson-Oppen procedure!

Non-deterministic Nelson-Oppen

- Deterministic Nelson-Oppen procedure doesn't work in this example because \mathcal{T}_{LIA} is **not convex**: $1 \leq x \leq 2 \models x = 1 \vee x = 2$ but in general neither $1 \leq x \leq 2 \not\models x = 1$ nor $1 \leq x \leq 2 \not\models x = 2$
- However, there is a **non-deterministic** Nelson-Oppen procedure that also works on non-convex theories
 - We still need **disjoint** and **stably-infinite** theories
- It works through **arrangements** of shared constants, basically doing case splitting $x = y \vee x \neq y$ between pair of shared constants x, y
 - Unsurprisingly, **exponential** worst-case time complexity

Optimization Modulo Theory

Extensions

- There are several **extensions** and **enhancements** to the SMT framework seen so far, e.g.
- Quantified formulas
- Layered solvers
- On-demand solvers
- **Optimization Modulo Theory**

Optimization Modulo Theory

- **OMT** is an extension of SMT where we need to find a model for an input formula φ that is **optimal** w.r.t. an **objective function** f_{obj}
- φ refers to a theory $\mathcal{T} = \mathcal{T}_{\preceq} \cup \mathcal{T}_1 \cup \dots \cup \mathcal{T}_n$ where
 - \mathcal{T}_{\preceq} contains a predicate \preceq representing a **total order**
 - $\bigcup_{i=1}^n \mathcal{T}_i$ might be empty
- The goal is finding a model \mathcal{M} s.t. $\varphi^{\mathcal{M}} = \text{true}$ and $f_{obj}^{\mathcal{M}}$ is minimal according to \preceq
 - Maximizing $f_{obj} \equiv$ minimizing $-f_{obj}$
- Typically, \preceq is the \leq predicate over integers or reals
 - E.g. $\mathcal{T}_{\mathcal{LIRA}} + \text{Nelson-Oppen } \mathcal{T}_i$

Optimization Modulo Theory

- OMT is “much younger” than SMT: first proposal in 2006
 - R. Nieuwenhuis and A. Oliveras. *On SAT Modulo Theories and Optimization Problems*. In SAT, volume 4121 of LNCS. Springer, 2006
- Nowadays different OMT proposals (see Sebastiani et al. works)
 - Max-SMT
 - Bit-vectors
 - Floating points
 - ...
- Some state-of-the-art SMT solvers natively provide OMT capabilities (Z3, OptiMathSAT) but others still don't (e.g. CVC5)
- Let's see an example by R. Sebastiani of $\text{OMT}(\mathcal{LRA})$ with linear search

OMT(\mathcal{LRA}) example


[w. pure-literal filt. \Rightarrow partial assignments]

- OMT(\mathcal{LRA}) problem:

$$\begin{aligned}\varphi \stackrel{\text{def}}{=} & (\neg A_1 \vee (2x + y \geq -2)) \\ & \wedge (A_1 \vee (x + y \geq 3)) \\ & \wedge (\neg A_2 \vee (4x - y \geq -4)) \\ & \wedge (A_2 \vee (2x - y \geq -6)) \\ & \wedge (\text{cost} < -0.2) \\ & \wedge (\text{cost} < -1.0) \\ & \wedge (\text{cost} < -2.0)\end{aligned}$$

$$\text{cost} \stackrel{\text{def}}{=} x$$

- $\mu = \left\{ \begin{array}{l} A_1, \neg A_1, A_2, \neg A_2, \\ (4x - y \geq -4), \\ (x + y \geq 3), \\ (2x + y \geq -2), \\ (2x - y \geq -6) \\ (\text{cost} < -0.2) \\ (\text{cost} < -1.0) \\ (\text{cost} < -2.0) \end{array} \right\}$



OMT(\mathcal{LRA}) example

[w. pure-literal filt. \implies partial assignments]


- OMT(\mathcal{LRA}) problem:

$$\begin{aligned}\varphi &\stackrel{\text{def}}{=} (\neg A_1 \vee (2x + y \geq -2)) \\ &\wedge (A_1 \vee (x + y \geq 3)) \\ &\wedge (\neg A_2 \vee (4x - y \geq -4)) \\ &\wedge (A_2 \vee (2x - y \geq -6)) \\ &\wedge (\text{cost} < -0.2) \\ &\wedge (\text{cost} < -1.0) \\ &\wedge (\text{cost} < -2.0)\end{aligned}$$

$$\text{cost} \stackrel{\text{def}}{=} x$$

$$\mu = \left\{ \begin{array}{l} A_1, \neg A_1, A_2, \neg A_2, \\ (4x - y \geq -4), \\ (x + y \geq 3), \\ (2x + y \geq -2), \\ (2x - y \geq -6) \\ (\text{cost} < -0.2) \\ (\text{cost} < -1.0) \\ (\text{cost} < -2.0) \end{array} \right\}$$

$\implies \text{SAT}, \min = -0.2$



OMT(\mathcal{LRA}) example

[w. pure-literal filt. \implies partial assignments]


- OMT(\mathcal{LRA}) problem:

$$\begin{aligned} \varphi &\stackrel{\text{def}}{=} (\neg A_1 \vee (2x + y \geq -2)) \\ &\wedge (A_1 \vee (x + y \geq 3)) \\ &\wedge (\neg A_2 \vee (4x - y \geq -4)) \\ &\wedge (A_2 \vee (2x - y \geq -6)) \\ &\wedge (\text{cost} < -0.2) \\ &\wedge (\text{cost} < -1.0) \\ &\wedge (\text{cost} < -2.0) \end{aligned}$$

$$\text{cost} \stackrel{\text{def}}{=} x$$

$$\mu = \left\{ \begin{array}{l} A_1, \neg A_1, A_2, \neg A_2, \\ (4x - y \geq -4), \\ (x + y \geq 3), \\ (2x + y \geq -2), \\ (2x - y \geq -6) \\ (\text{cost} < -0.2) \\ (\text{cost} < -1.0) \\ (\text{cost} < -2.0) \end{array} \right\}$$

$\implies \text{SAT}, \min = -1.0$



OMT(\mathcal{LRA}) example

[w. pure-literal filt. \Rightarrow partial assignments]


- OMT(\mathcal{LRA}) problem:

$$\begin{aligned}\varphi &\stackrel{\text{def}}{=} (\neg A_1 \vee (2x + y \geq -2)) \\ &\wedge (A_1 \vee (x + y \geq 3)) \\ &\wedge (\neg A_2 \vee (4x - y \geq -4)) \\ &\wedge (A_2 \vee (2x - y \geq -6)) \\ &\wedge (\text{cost} < -0.2) \\ &\wedge (\text{cost} < -1.0) \\ &\wedge (\text{cost} < -2.0)\end{aligned}$$

$$\text{cost} \stackrel{\text{def}}{=} x$$

$$\mu = \left\{ \begin{array}{l} A_1, \neg A_1, A_2, \neg A_2, \\ (4x - y \geq -4), \\ (x + y \geq 3), \\ (2x + y \geq -2), \\ (2x - y \geq -6) \\ (\text{cost} < -0.2) \\ (\text{cost} < -1.0) \\ (\text{cost} < -2.0) \end{array} \right\}$$

$\Rightarrow \text{SAT}, \min = -2.0$



OMT(\mathcal{LRA}) example

[w. pure-literal filt. \Rightarrow partial assignments]


• OMT(\mathcal{LRA}) problem:

$$\begin{aligned} \varphi &\stackrel{\text{def}}{=} (\neg A_1 \vee (2x + y \geq -2)) \\ &\wedge (A_1 \vee (x + y \geq 3)) \\ &\wedge (\neg A_2 \vee (4x - y \geq -4)) \\ &\wedge (A_2 \vee (2x - y \geq -6)) \\ &\wedge (\text{cost} < -0.2) \\ &\wedge (\text{cost} < -1.0) \\ &\wedge (\text{cost} < -2.0) \end{aligned}$$

$$\text{cost} \stackrel{\text{def}}{=} x$$

$$\bullet \mu = \left\{ \begin{array}{l} A_1, \neg A_1, A_2, \neg A_2, \\ (4x - y \geq -4), \\ (x + y \geq 3), \\ (2x + y \geq -2), \\ (2x - y \geq -6) \\ (\text{cost} < -0.2) \\ (\text{cost} < -1.0) \\ (\text{cost} < -2.0) \end{array} \right\}$$

\Rightarrow UNSAT, $\min = -2.0$



Offline OMT(\mathcal{LRA})

- **Linear search** repeatedly narrows the cost domain $[l_i, u_i)$ by adding $cost < c_i$ if a model with cost c_i is found at the i -th iteration
 - If no model is found, c_i is the minimum cost
- **Binary search** picks a **pivot** $p_i \in [l_i, u_i)$ and adds $cost < p_i$
 - $p_i \simeq (l_i + u_i)/2$
 - If no model is found, look into $[p_i, u_i)$
 - Can be more efficient, but we **must** know the cost **bounds**
- This approach is called **offline** because the SMT solvers used to find the models are **black-boxes**
 - No need to change their internals

Offline OMT(\mathcal{LRA})

Algorithm 1 Offline OMT($\mathcal{LA}(\mathbb{Q})$) Procedure based on Mixed Linear/Binary Search.

Require: $\langle \varphi, \text{cost}, \text{lb}, \text{ub} \rangle$ {ub can be $+\infty$, lb can be $-\infty$ }

```
1:  $l \leftarrow \text{lb}; u \leftarrow \text{ub}; \text{PIV} \leftarrow \top; \mathcal{M} \leftarrow \emptyset$ 
2:  $\varphi \leftarrow \varphi \cup \{\neg(\text{cost} < l), (\text{cost} < u)\}$ 
3: while ( $l < u$ ) do
4:   if (BinSearchMode()) then {Binary-search Mode}
5:      $\text{pivot} \leftarrow \text{ComputePivot}(l, u)$ 
6:      $\text{PIV} \leftarrow (\text{cost} < \text{pivot})$ 
7:      $\varphi \leftarrow \varphi \cup \{\text{PIV}\}$ 
8:      $\langle \text{res}, \mu \rangle \leftarrow \text{SMT.IncrementalSolve}(\varphi)$ 
9:      $\eta \leftarrow \text{SMT.ExtractUnsatCore}(\varphi)$ 
10:   else {Linear-search Mode}
11:      $\langle \text{res}, \mu \rangle \leftarrow \text{SMT.IncrementalSolve}(\varphi)$ 
12:      $\eta \leftarrow \emptyset$ 
13:   end if
14:   if ( $\text{res} = \text{SAT}$ ) then
15:      $\langle \mathcal{M}, u \rangle \leftarrow \text{Minimize}(\text{cost}, \mu)$ 
16:      $\varphi \leftarrow \varphi \cup \{(\text{cost} < u)\}$ 
17:   else { $\text{res} = \text{UNSAT}$ }
18:     if ( $\text{PIV} \notin \eta$ ) then
19:        $l \leftarrow u$ 
20:     else
21:        $l \leftarrow \text{pivot}$ 
22:        $\varphi \leftarrow \varphi \setminus \{\text{PIV}\}$ 
23:        $\varphi \leftarrow \varphi \cup \{\neg \text{PIV}\}$ 
24:     end if
25:   end if
26: end while
27: return  $\langle \mathcal{M}, u \rangle$ 
```

Annotations:

- Lines 14-16: $u = \text{current best bound}$
- Lines 17-19: Linear search completed
- Lines 20-23: Updating binary search pivot

From R. Sebastiani, S. Tomasi: *Optimization Modulo Theories with Linear Rational Costs*. ACM Trans. Comput. Log. 16(2): 12:1-12:43 (2015)

Inline OMT(\mathcal{LRA})

- The minimal cost is computed by a **minimizer** over linear rational inequalities
 - E.g., standard **simplex** techniques
- The offline approach can be improved by an **inline** schema
 - More **efficient**, but it requires **modifying the internals** of SMT solver
- In a nutshell, the inline approach **integrates** the optimization procedure **into** the SMT solver

Take-home messages

- Different **theory solvers** have been developed for different theories
 - E.g. EUF, DL, LRA, LIA, ...
- We often need to **combine** theories
 - Under certain conditions, **Nelson-Oppen** procedure can be used
- SMT solving can be optimized and **extended**
 - **Optimization modulo theory**

- Handbook of Satisfiability – Chapter 12 “*Satisfiability Modulo Theories*” by C. Barrett, R. Sebastiani, S.A. Seshia, C. Tinelli
 - Search “Satisfiability Modulo Theories - EECS at UC Berkeley”
- Barrett, Clark, and Cesare Tinelli. “Satisfiability modulo theories.” Handbook of model checking. Springer, Cham, 2018. 305-343.
- SAT/SMT schools
 - <https://sat-smt.in/>
- ...