

6. (Mixed) Integer Linear Programming

Roberto Amadini

Department of Computer Science and Engineering, University of Bologna, Italy

Combinatorial Decision Making and Optimization

2nd cycle degree programme in Artificial Intelligence

University of Bologna, Academic Year 2024/25



From Reals to Integers

- We have seen the **simplex method** to solve LP problems in “*typically polynomial*” time
- We however can **always** solve LP problems in **polynomial time**
 - **Ellipsoid** methods (inefficient, worse than simplex in practice)
 - **Interior point** methods (e.g., Karmarkar, can outperform simplex)
- If the **domain** of the variables involved is \mathbb{R} or \mathbb{Q} solving LP problems is **not NP-hard**
- What if we require domains to be subsets of \mathbb{Z} or \mathbb{N} ? Would the problem be easier? Or harder? Or the same?
 - Let's go back to the baking example

Baking cakes

```
1 % We know how to make two sorts of cakes. A banana cake which takes 250g of
2 % self-raising flour, 2 mashed bananas, 75g sugar and 100g of butter, and a
3 % chocolate cake which takes 200g of self-raising flour, 75g of cocoa, 150g
4 % sugar and 150g of butter. We can sell a chocolate cake for $4.50 and a
5 % banana cake for $4.00. And we have 4kg self-raising flour, 6 bananas,
6 % 2kg of sugar, 500g of butter and 500g of cocoa. How many of each sort of
7 % cake should we bake for the fete to maximise the profit
8
9 var 0..100: b; % no. of banana cakes
10 var 0..100: c; % no. of chocolate cakes
11
12 % flour
13 constraint 250*b + 200*c <= 4000;
14 % bananas
15 constraint 2*b <= 6;
16 % sugar
17 constraint 75*b + 150*c <= 2000;
18 % butter
19 constraint 100*b + 150*c <= 500;
20 % cocoa
21 constraint 75*c <= 500;
22
23 % maximize our profit
24 solve maximize 400*b + 450*c;
25
26 output ["no. of banana cakes = \"(b)\\n\",
27         "no. of chocolate cakes = \"(c)\\n\""];
```

Baking cakes




- $b, c \in 0..100 = \{0, 1, \dots, 100\} \subsetneq [0, 100] = \{x \in \mathbb{R} \mid 0 \leq x \leq 100\}$
- $b = 3, c = 4/3$ can't be a feasible solution anymore!

Baking cakes

- On the previous episodes, we assumed we could sell **slices** of cakes
 - E.g., one third of banana cake and two fifths of chocolate cake
- This could be a **too strong** assumption in a **real-world** setting
 - Imagine we are selling cars instead of cakes
- The **feasible/optimal regions** are now **sets** of integer points
 - “Grids” instead of polyhedra
- What is the **optimal** solution if we strictly require $b, c \in 0..100$?

Baking cakes



- $b, c \in 0..100 = \{0, 1, \dots, 100\} \subsetneq [0, 100] = \{x \in \mathbb{R} \mid 0 \leq x \leq 100\}$
- Optimal solution $b = c = 2$


Integer Linear Programming

- An **Integer Linear Programming (ILP)** problem has **standard** form:

$$\begin{array}{ll}\max & \sum_{j=1}^n c_j x_j \\ \text{s.t.} & \sum_{j=1}^n a_{i,j} x_j = b_i \quad 1 \leq i \leq m \\ & x_j \geq 0, \quad x_j \in \mathbb{Z} \quad 1 \leq j \leq n\end{array}$$

- If we require only $k < n$ variables to be integers, then it is a **Mixed-Integer Programming (MIP)** problem
 - ILP a.k.a. Integer Programming (**IP**)
 - MIP a.k.a. Mixed-Integer Linear Programming (**MILP**)
- **Linear relaxation**: MIP problem with **no integrality constraints** $x_j \in \mathbb{Z}$
 - If $\mathcal{L}(P)$ is the linear relaxation of P , $\mathcal{F}_P \subseteq \mathcal{F}_{\mathcal{L}(P)}$ hence **solving** $\mathcal{F}_{\mathcal{L}(P)}$ and **rounding** its optimal solution does **not work** in general

Linear relaxation



- Optimal solution of linear relaxation: $b = 3$, $c = 4/3 = 1.\bar{3}$, rounding to nearest integer $b = 3$, $c = 1$. Obj. value = $400 \cdot 3 + 450 \cdot 1 = 1650$
- Optimal solution original problem: $b = 2$, $c = 2$. Obj. value = $400 \cdot 2 + 450 \cdot 2 = 1700$

Linear relaxation

- Rounding the optimal solution of $\mathcal{L}(P)$ is **not sound** in general!
- E.g. with $P : \max(x)$ s.t. $x \leq 5/3, x \geq 0, x \in \mathbb{Z}$ we have that $5/3 = 1.\bar{6}$ is optimal for $\mathcal{L}(P)$ but its rounding is $2 \notin \mathcal{F}_P = \{0, 1\}$
 - In this case $\mathcal{O}_P = \{1\} \neq \{5/3\} = \mathcal{O}_{\mathcal{L}(P)}$
- Clearly, $\mathcal{F}_{\mathcal{L}(P)} = \emptyset \implies \mathcal{F}_P = \emptyset$: solving $\mathcal{F}_{\mathcal{L}(P)}$ can help to detect **unsatisfiability** of P
- If $\mathcal{L}(P)$ **unbounded**, P can be:
 - bounded, e.g. $P : \max(x_1)$ s.t. $x_1 = \sqrt{2}x_2$ and $x_1, x_2 \in \mathbb{N}$
 - unbounded, e.g. $P : \max(x)$ s.t. $x \in \mathbb{N}$
 - unsatisfiable, e.g. $P : \max(x_1)$ s.t. $0.1 \leq x_2 \leq 0.2$ and $x_1, x_2 \in \mathbb{N}$

Complexity

- Adding integrality has huge impact on the **complexity** of LP solving
- **No known** algorithms for solving MIP problems in **polynomial time**
 - otherwise, **P = NP**
- More precisely, MIP problems are **NP-complete**
 - They are in **NP** (**certifying** solutions takes **polynomial** time)
 - *Solvable by NDTM in polynomial time...*
 - They are among the *hardest* problems in **NP** (**NP-hard**)
- If we could solve a generic MIP problem in polynomial time, we could also solve, e.g., **any SAT problem** in polynomial time
 - And **all** the NP problems in polynomial time...

SAT to MIP reduction

- From any **generic SAT** problem P_{SAT} with clauses C_1, \dots, C_m and literals ℓ_1, \dots, ℓ_n we get an **equisatisfiable MIP** problem P_{MIP} with n variables x_j (one per literal) and m constraints (one per clause):

$$\begin{array}{ll}\max & 0 \\ \text{s.t.} & \sum_{\ell_j \in P_i} x_j + \sum_{\ell_j \in N_i} (1 - x_j) \geq 1 \quad 1 \leq i \leq m \\ & x_j \in \{0, 1\} \quad 1 \leq j \leq n\end{array}$$

where P_i and N_i are the **positive** and **negative** literals of clause C_i

- E.g. $\ell_1 \vee \neg \ell_2 \vee \ell_4 \vee \neg \ell_5 \implies x_1 + (1 - x_2) + x_4 + (1 - x_5) \geq 1 \implies x_1 - x_2 + x_4 - x_5 \geq -1 \implies -x_1 + x_2 - x_4 + x_5 \leq 1$

SAT to MIP reduction

- Reducing P_{SAT} to P_{MIP} takes **polynomial** time
 - $O(mn)$ time
- P_{SAT} feasible $\iff P_{MIP}$ feasible, and any solution of P_{MIP} can be mapped back into a solution of P_{SAT} in **polynomial** time
 - $O(n)$ time: $x_j = 0 \mapsto \ell_j = \text{false}$, $x_j = 1 \mapsto \ell_j = \text{true}$
- If P_{MIP} solvable in polynomial time \implies **any SAT** problem solvable in poly. time \implies **any NP-complete** problem in poly. time \implies **P=NP**
- Because rounding is in general not applicable, we have to tackle MIP solving with other approaches

Handling NP-hardness

Different ways of handling NP-hard problems:

- **Exact** algorithms: they guarantee to find an **optimal** solution although this may take **exponential** time
 - The focus of **this** course
- **Approximation** algorithms: they guarantee in **polynomial time** a **(sub-)optimal** solution at most ρ times worse the optimal one
 - $\rho =$ **approximation factor**
- **Heuristic** algorithms: **no guarantee** of optimality nor polynomial runtime, but “in practice” they find **good** solutions in **reasonable** time
 - According to **empirical** evidences

Branch-and-bound

- **Branch-and-bound (BB)** is based on the **divide-et-impera** principle: split a “big” problem into **sub-problems** until a success or a failure
 - **Overall solution** derived from the solution of sub-problems
- **Branch** phase = explore the sub-problems
 - sub-trees of the **search tree**
- **Bound** phase = compute the **bounds** of sub-problem optimal solution
 - to possibly **prune** the search tree if current solution not improvable
- BB is a **general paradigm** applicable to various NP-hard problems

Branch-and-bound


- We can solve a standard MIP problem P via BB. Suppose $P_0 = \mathcal{L}(P)$ has solution β_1, \dots, β_n with some $\beta_k \notin \mathbb{Z}$ (very lucky otherwise!)
- Pick a x_k s.t. $\beta_k \notin \mathbb{Z}$ and branch:
$$\begin{cases} P_1 = P_0 \cup \{x_k \leq \lfloor \beta_k \rfloor\} \\ P_2 = P_0 \cup \{x_k \geq \lceil \beta_k \rceil\} \end{cases}$$
 - Note $\mathcal{F}_P = \mathcal{F}_{P_1} \cup \mathcal{F}_{P_2}$ and $\mathcal{F}_{P_1} \cap \mathcal{F}_{P_2} = \emptyset$
- We can solve P_1, P_2 to optimality and take the best solution (if any)
 - If integral, optimal for P too! Otherwise, branch again on P_1, P_2, \dots
- So we build a search tree rooted in P_0 with edges $P_i \rightarrow P_j$ if child node P_j is a sub-problem of parent node P_i

Branch-and-bound

- If a P_k has **integral** optimal solution, compare its obj. value z_k with the best obj. value z^* so far (**best bound**): if $z_k > z^*$, then $z^* \leftarrow z_k$
- Otherwise, we **cannot improve** z^*
 - In any case, **fathom** this node: P_k will be a **leaf**
 - Initially, $z^* \leftarrow -\infty$
- If P_k is **unfeasible** or has **non-integral** solution with obj. value $z_k \leq z^*$ also fathom its node (**prune** sub-tree rooted in P_k)
- At the end, the optimal solution for P is a **leaf** P_k with obj. value z^*
 - Leaves also called **fathomed nodes**
 - Current optimal solution also called **incumbent solution**

Baking example


Let's how BB works on the baking problem where $B, C \in 0..100$




If we solve $P_0 = \mathcal{L}(P)$, optimal solution $B = 3$, $C = 4/3$ not feasible: we need to **branch** on C :

$$\begin{cases} P_1 = P_0 \cup \{C \leq \lfloor 4/3 \rfloor = 1\} \\ P_2 = P_0 \cup \{C \geq \lceil 4/3 \rceil = 2\} \end{cases}$$

Baking example



The resulting **search tree** is:




remember objective function is $Z = 400B + 450C$

Baking example




If $C \leq 1$, optimal solution is integral: $B = 3, C = 1$ with value 1650




Node P_1 is a **leaf**: we **backtrack** and explore P_2

Baking example



If $C \geq 2$, we get a better solution $B = 2, C = 2$ with value 1700



search completed: solution of node P_2 is optimal

Branch and Bound

- BB works well typically in **combination** with other techniques
 - Presolve
 - Cutting planes
 - Heuristics
 - Parallelism
 - ...


- **Presolve** means **reformulating** a problem **before** its actual solving process to possibly reduce its size
 - Presolve should be computationally **efficient**
- **Bounds tightening**, e.g.: $\{x_1 + x_2 \geq 20, x_1 \leq 10\} \models x_2 \geq 10$
 - If $x_2 \in 1..9$, we **eagerly** detect unsatisfiability
- **Problem reduction**, e.g.: $\{x_1 + x_2 \leq 0.8\} \models x_1 = x_2 = 0$
 - x_1, x_2 can be **removed** from problem formulation
- Pre-processing a MIP problem P is important because it can **reduce** the size of $\mathcal{F}_{\mathcal{L}(P)}$ **without altering** \mathcal{F}_P

Cutting planes

- Cutting planes allowed for significant advancements of MIP solving
- Instead of (*in addition to...*) branching on sub-problems, we repeatedly add new constraints entailed by the original problem P
- The idea is to “cut out” parts of $\mathcal{F}_{\mathcal{L}(P)} - \mathcal{F}_P$ along the solving process to remove non-integral solutions
 - Until we converge to an optimal solution for P
- The existence of a cut separating optimal solution in $\mathcal{F}_{\mathcal{L}(P)} - \mathcal{F}_P$ from $\mathcal{F}(P)$ is always guaranteed
 - But not its uniqueness!

Cutting planes

Formally, a **cut** for a MIP problem P in standard form is an inequality $px \leq q$ such that $py \leq q$ and $pz > q$ for each $y \in \mathcal{F}_P$ and $z \in \mathcal{O}_{\mathcal{L}(P)}$



Cutting Planes

Gomory's cut

- Proposed by R. Gomory in the 1950s
- Suppose basis $\mathcal{B}^* = \{x_{i_1}, \dots, x_{i_m}\}$ is optimal for $\mathcal{L}(P)$. The problem in terms of the basic variables (scalar form) is:

$$x_{i_k} = \beta_k + \sum_{j=1}^{n-m} \alpha_{k,j} x_{i_{m+j}} \quad \text{for } k = 1, \dots, m$$

where $x_{i_{m+1}}, x_{i_{m+2}}, \dots, x_{i_n}$ are non-basic variables

- The optimal solution x^* is $x_{i_k}^* = \beta_k$ for $k = 1, \dots, m$ and $x_{i_j}^* = 0$ for $j = m+1, \dots, n$. If there is a $\beta_k \notin \mathbb{Z}$, then $x^* \in \mathcal{F}_{\mathcal{L}(P)} - \mathcal{F}_P$ so we generate a cut to separate x^* from \mathcal{F}_P

Gomory's cut

- Gomory's cut has the form $-f_k + \sum_{j=1}^{n-m} f_{k,j} x_{i_{m+j}} \geq 0$ where:
 - $f_k = \beta_k - \lfloor \beta_k \rfloor$
 - $f_{k,j} = -\alpha_{k,j} - \lfloor -\alpha_{k,j} \rfloor$
 - $0 < f_k < 1$ and $0 \leq f_{k,j} < 1$ by definition of $\lfloor \cdot \rfloor$ and because $\beta_k \notin \mathbb{Z}$
 - So $-f_k + \sum_{j=1}^{n-m} f_{k,j} x_{i_{m+j}}^* < 0$: the cut is **violated** by optimal solution x^*
 - We can prove that $-f_k + \sum_{j=1}^{n-m} f_{k,j} x_{i_{m+j}} \geq 0$ holds **for each** $x \in \mathcal{F}_P$
- Then $\mathcal{L}(P) \leftarrow \mathcal{L}(P) \cup \left\{ y_k = -f_k + \sum_{j=1}^{n-m} f_{k,j} x_{i_{m+j}}, \quad y_k \geq 0 \right\}$
 - A new **slack variable** y_k added at each round
- $\mathcal{B}' \leftarrow \mathcal{B}^* \cup \{y_k\}$ and solve with **dual simplex**
 - \mathcal{B}' now **unfeasible** ($y_k = -f_k < 0$) but **dual feasible**

Example

- E.g., for baking example we get $B^* = \{S_1, B, S_3, C, S_5\}$, with optimal solution $B = 3$, $C = 4/3$, $S_1 = \dots = S_5 = 0$ having value $z^* = 1800$
 - $\alpha_{C,S_2} = 1/3$, $\alpha_{C,S_4} = 1/150$
- $f_C = \frac{4}{3} - \lfloor \frac{4}{3} \rfloor = \frac{4}{3} - 1 = \frac{1}{3}$
- $f_{C,S_2} = -\frac{1}{3} - \lfloor -\frac{1}{3} \rfloor = -\frac{1}{3} + 1 = \frac{2}{3}$
- $f_{C,S_4} = -\frac{1}{150} - \lfloor -\frac{1}{150} \rfloor = \frac{1}{150} + 1 = \frac{149}{150}$
- Gomory's cut is $y_C = -\frac{1}{3} + \frac{2}{3} \cdot S_2 + \frac{149}{150} \cdot S_4 \geq 0$
 - This constraint "cuts-out" current non-integral solution
- Starting basis for dual simplex is $\{S_1, B, S_3, C, S_5, y_C\}$, then y_C out, choose entering variable, reformulate with new basis...

Branch-and-cut

- Gomory's cut considered **ineffective at the time** because of numerical instability and number of cuts needed for convergence
- In mid 1990s, **G. Cornuéjols** et al. proved it to be very effective in combination with branch-and-bound: **branch-and-cut**
- Basically, it runs **BB** for P and if an optimal solution of $\mathcal{L}(P)$ is not integral it adds cutting planes to **tighten $\mathcal{L}(P)$**
 - E.g., <https://www.ibm.com/docs/en/icos/22.1.0?topic=c-branch-cut-in-cplex-1>
- Different cutting planes algorithms (or **separators**) can be used
 - ...and **learned**: <https://arxiv.org/abs/2311.05650>

Row and Column Generation

- Cutting planes can be seen as “row generation” methods: new constraints are added at each step
- Bender's decomposition is another row generation method dividing a problem P into master problem (min) and sub-problem(s) (max)
 - Idea: iteratively fix the value of some variables and solve the dual of residual sub-problem to get cuts or better objective bounds
 - Logic-Based Bender's Decomposition: sub-problems are generic problems solvable with any solver (e.g., CP or SMT solvers)
- Also column-generation methods exist: start with subset of variables and repeatedly add variables until objective value cannot be improved
 - Assumption: only a small subset of variables is useful
 - E.g., Dantzig-Wolfe decomposition

Heuristics

- **Heuristics** methods aim to find “good” solutions in “reasonable” time
 - Inherently **empirical**, weak theoretical guarantees
- **Constructive** methods: start with empty solution and iteratively **extend** it to a complete solution
- **Local search** methods: start with a complete solution and iteratively **modify** parts to improve it
- **Evolutionary** methods: explore a **population** of solutions using mutation, crossover, and selection
- **Meta-heuristics** methods: higher-level strategies for selecting, combining, tuning, or guiding other heuristics

Primal heuristics

- Modern MIP solvers often use so-called **primal heuristics** in addition to branching and cutting, e.g.:
 - **Rounding** non-integral solutions and check
 - Easy and fast, typically used at root node
 - **Diving**: iteratively fix fractional variables to some value and solve the remaining sub-problem
 - Helps to quickly find feasible solutions

Primal heuristics

- **Neighborhood-based**: improve incumbent by solving sub-problems “around it” by fixing some incumbent variables
 - Relaxation-Induced Neighborhood Search (RINS), Large Neighborhood Search (LNS), Local branching, ...
- **Node Selection**: decide specific nodes to apply heuristics selectively rather than at every node
 - E.g. according to depth, variable fixings, dual gap, ...
- **Adaptive** dynamic heuristic selection and configuration of heuristics
 - Parameters tuning, portfolios, ML methods...

Warm starts

- MIP solving can be **warm-started** with initial value(s) to (some of) the variables
- This **not necessarily** translates to a new **incumbent**: warm-start solution might be unfeasible or not better than current best bound
 - Or, if partial, it might take too long to compute a complete solution
- A warm start vector might come from the **knowledge** of a similar problem, or from the **expertise** of a domain expert
- E.g., warm-start the MIP solver by reusing a route plan or a timetable from a previous day or period

Take-home messages

- Adding **integer variables** to LP significantly increases its complexity
- **MIP solving** can be tackled with (a combination of) different approaches
 - Exact
 - Approximate
 - Heuristic
- **Rounding** non-integral solutions of **linear relaxation** does **not work** in general

Take-home messages

- **Branch-and-bound**: divide-et-impera approach, **branches** on variables with non-integer value, stop if we cannot improve **incumbent** solution
- **Cutting planes**: linear equalities **separating** non-integral, optimal solutions of linear relaxation from feasible region of original problem
 - Gomory's cut
 - Branch-and-cut
- **Heuristics**: no strong guarantees on optimality/runtime, but “in practice” they find **good** solutions in **reasonable** time
 - Inherently **empirical**