

# Relevant Architectures (1)




- Encoder-classifier
  - Transfer Learning
- Encoder-decoder
- Autoencoder
- U-net
- Transformer (next lesson)
- Visual Transformer (next lesson)

# Encoder-classifier



AlexNet Architecture (Krizhevsky, Sutskever e Hinton), winner of a NIPS contest in 2012.




Purely historical interest.

In deep convolutional networks, it is common practice to alternate convolutional layers with **pooling** layers, where each neuron simply takes the mean or maximal value in its receptive field.

This has a double advantage:

- ▶ it reduces the dimension of the output
- ▶ it gives some tolerance to translations


VGG 16 (Simonyan e Zisserman). 92.7 accuracy (top-5) in ImageNet (14 millions images, 1000 categories).



Picture by Davi Frossard: VGG in TensorFlow

# Inception V3

## Inception V3




The convolutional part is a long composition of  
inception modules



# Inception modules

The networks is composed of inception modules (towers of nets):



Video from the Udacity course "Deep Learning"






# Variants

The point is to induce the net to learn different filters.

Many variants proposed and used over years:



# Xception, MobileNet, EfficientNet


---

These models focus on reducing the number of parameters for efficiency reasons (e.g. embedding on mobile devices)

- **Xception**
- **MobileNet** - a class of “light” models conceived to be deployed on mobile devices.
- **EfficientNet**: Rethinking Model Scaling for Convolutional Neural Networks.


One of the key tools are **Depth Separable Convolutions**.

# Efficient Net





# Residual Learning

Another recent topic is residual learning.




Instead of learning a function  $\mathcal{F}(\mathbf{x})$  you try to learn  $\mathcal{F}(\mathbf{x}) + \mathbf{x}$ .

# The right intuition



# Residual networks



you add a residual shortcut connection every 2-3 layers

Inception Resnet is an example of a such an architecture

# Why Residual Learning works?

---


- ▶ it seems to be a good idea to try to learn **non-linear corrections** over a linear baseline
- ▶ during back propagation, **the gradient at higher layers can easily pass to lower layers**, without being mediated by the weight layers, which may cause vanishing gradient or exploding gradient problem.

[Demo]





# A parenthesis: Transfer Learning



We learned that the first layers of convolutional networks for computer vision compute feature maps of the original image of growing complexity.

The filters that have been learned (in particular, the most primitive ones) are likely to be **independent from the particular kind of images they have been trained on.**

They have been trained on a **huge amount of data** and are probably very good.


It is a good idea to try to *reuse them* for other classification tasks.

# Transfer Learning with CNNs

## Transfer Learning with CNNs



1. Train on ImageNet



2. If small dataset: fix all weights (treat CNN as fixed feature extractor), retrain only the classifier

i.e. swap the Softmax layer at the end



3. If you have medium sized dataset, “**finetune**” instead: use the old weights as initialization, train the full network or only some of the higher layers

retrain bigger portion of the network, or even all of it.

# When Transfer Learning makes sense


---

transferring knowledge from problem A to problem B makes sense if

- the two problems have “similar” inputs
- we have much more training data for A than for B

# What we may expect

## Faster and more accurate training



# Encoder-Decoder




Encoder-Decoder models are a large class of models composed of two basic components:

**Encoder** transforming the input into a latent representation, typically a vector in some space, called the **latent space**.

**Decoder** reconstructing from the internal representation an explicit output into the visible space.

# Autoencoders

An autoencoder is a net trained to reconstruct **input data** out of a learned internal representation



Usually, the internal representation has lower dimensionality w.r.t. the input.



# Demo: autoencoders in Keras

---



See [Building autoencoders in Keras](#), on the Keras blog



Why is data compression possible, in general?

Because we exploit **regularities** (correlations) in the features describing input data.

If the input has a random structure (high **entropy**) no compression is possible

- random, lawless, uncompressible, high entropy
- ordered, lawfull, compressible, low entropy

Why is data compression possible, in general?

Because we exploit **regularities** (correlations) in the features describing input data.

If the input has a random structure (high **entropy**) no compression is possible

- random, lawless, uncompressible, high entropy
- ordered, lawfull, compressible, low entropy

# Autoencoder's compression

---

When the internal layer has fewer units of the input, autoencoders implement as a form **data compression**, similar to not-linear PCA.

- ▶ **data-specific**: it only works well on data with strong correlations (e.g. digits, faces, etc.) This is different from traditional data compression algorithms
- ▶ **lossy**: the output is degraded with respect to the input. This is different from textual compression algorithms, such as gzip
- ▶ **directly trained** on unlabeled data samples. We usually talk of **self-supervised** training



# Anomaly detection

Autoencoding is **data specific**: the autoencoder works well on data similar to those it was trained on.

If applied on different data (anomaly), it will perform poorly.

Example on mnist data with the autoencoder of the previous demo.  
mean loss = 0.10, std: 0.03



loss = 0.14 OK!



loss = 0.27 LIAR!



# U-net

Suggested reading:

U-Net: Convolutional Networks for Biomedical Image Segmentation. By  
O.Ronneberger, P.Fischer, T.Brox

The U-net is the de-facto standard for image-to-image processing.

Typical applications comprise

- ▶ segmentation
- ▶ denoising
- ▶ deblurring
- ▶ image restoration
- ▶ inpainting
- ▶ ...



# U-net architecture

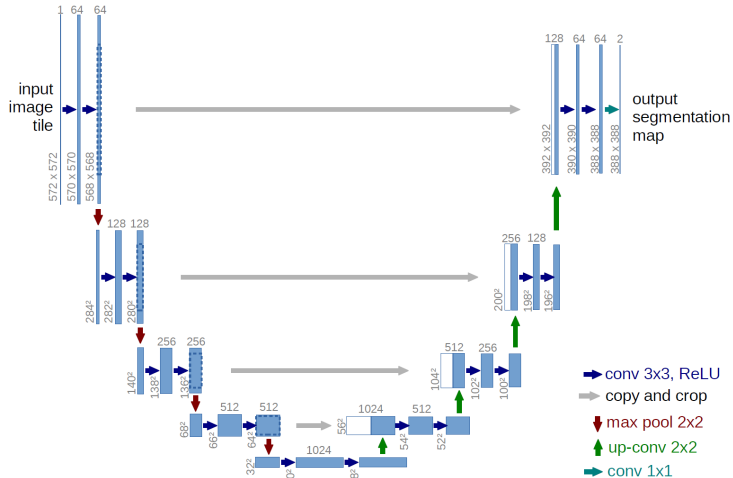
---

The U-net complements an autoencoder structure with **skip connections** concatenating layers of the encoder to layers of the decoder at similar spatial resolution.

The autotencoder allows to obtain a holistic interpretation of the input (the what), while skip connections allows a more precise location and processing of features (the where).

The U-net allows to process together global features (extracted through the encoding process) with local features (merged with skip-connection).

# U-net architecture



[Demo]

